

# EDA

March 28, 2024

## 1 EDA - Spotify Top Songs and Audio Features

The goals of this analysis are: \* Loading the data into a code environment and in a dataframe that can be manipulated \* Getting familiar with the data \* Basic preprocessing, checking for anomalous data or wrongly formatted data or NaNs etc. \* First visualizations of the data and its distributions. \* Creating a new dataset that is indexed by artists and not by tracks and finding suitable ways to “average” the values across all the tracks from an artist.

Imports

```
[18]: #Imports
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from scipy.stats import mode
```

```
[19]: # Load the dataset into a DataFrame
df = pd.read_csv('data.csv')

# Display the first few rows of the DataFrame
df.head(10)
```

```
[19]:
```

	id	artist_names \
0	000xQL6tZNLJzIrtIgxqS1	ZAYN, PARTYNEXTDOOR
1	003eoIwxETJujVWmNFMoZy	Alessia Cara
2	003vvx7Niy0yvvhvHt4a68B	The Killers
3	00B7TZ0Xawar6NZ00JFomN	Cardi B, Chance the Rapper
4	00Blm7zeNqgYLPtW6zg8cj	Post Malone, The Weeknd
5	00EPIEnX1JFjff8sC6bccd	Thalia, NATTI NATASHA
6	00ETaeHUQ61ops3oWU1Wrt	Kygo, Donna Summer
7	00ZKeP47bZtswtANKvxz2j	Tropa do Bruxo, DJ Ws da Igrejinha, SMU, Triz,...
8	00gpGR84M27moP7AFuqHIX	YBN Nahmir
9	00imgaPLYRrMGn9o83hfmk	Brent Faiyaz

  

	track_name \
0	Still Got Time (feat. PARTYNEXTDOOR)
1	Growing Pains

2 Mr. Brightside  
3 Best Life (feat. Chance The Rapper)  
4 One Right Now (with The Weeknd)  
5 No Me Acuerdo  
6 Hot Stuff  
7 Baile do Bruxo  
8 Bounce Out With That  
9 LOOSE CHANGE

	source	key	mode	time_signature	\
0	RCA Records Label	G	Major	4 beats	
1	Def Jam Recordings	C#/Db	Minor	4 beats	
2	Island Records	C#/Db	Major	4 beats	
3	Atlantic/KSR	A	Major	4 beats	
4	Republic Records	C#/Db	Major	4 beats	
5	Sony Music Latin	G	Minor	4 beats	
6	RCA Records Label	F	Major	4 beats	
7	Tropa Do Bruxo	G	Minor	5 beats	
8	2018	G#/Ab	Major	4 beats	
9	Lost Kids LLC., Marketed by Venice / Stem	C#/Db	Minor	4 beats	

	danceability	energy	speechiness	acousticness	instrumentalness	\
0	0.748	0.627	0.0639	0.13100	0.000000	
1	0.353	0.755	0.7330	0.08220	0.000000	
2	0.352	0.911	0.0747	0.00121	0.000000	
3	0.620	0.625	0.5530	0.28700	0.000000	
4	0.687	0.781	0.0530	0.03610	0.000000	
5	0.836	0.799	0.0873	0.18700	0.000000	
6	0.681	0.773	0.1480	0.01900	0.000001	
7	0.734	0.228	0.5300	0.88900	0.006420	
8	0.857	0.560	0.1730	0.04260	0.000000	
9	0.574	0.369	0.0814	0.75300	0.000000	

	liveness	valence	loudness	tempo	duration_ms	weeks_on_chart	\
0	0.0852	0.524	-6.029	120.963	188491	17	
1	0.3900	0.437	-6.276	191.153	193680	2	
2	0.0995	0.236	-5.230	148.033	222973	125	
3	0.3140	0.665	-7.438	167.911	284856	2	
4	0.0755	0.688	-4.806	97.014	193507	30	
5	0.0920	0.772	-4.247	94.033	217653	16	
6	0.1100	0.429	-5.749	119.961	199008	1	
7	0.1020	0.522	-4.731	162.524	221538	3	
8	0.1530	0.482	-8.278	94.949	91011	6	
9	0.1470	0.440	-8.931	84.975	226011	1	

streams  
0 107527761

```
1    9944865
2    512388123
3    11985346
4    301860377
5    98123727
6    4569978
7    27916960
8    4913180
9    5854629
```

### 1.0.1 Inspecting the data

We notice that we have some basic information about each songs, like for example the artist name, the track name, the record label, these identify a song and is information that is searched by the user. We also have other metrics like the mode (Major or Minor) and the key, time signature, duration of the song and tempo. These are the musical “stats” of a track. In addition we have some spotify platform specific data, the number of streams for each track, and the number of weeks that the data was on charts. What makes the dataset very interesting to us however, is the addition of these tasteful features about each track such as “danceability”, “energy” and “valence” etc. These features are calculated by the author of the dataset, the code for the creation of the dataset can be viewed here: [github link to the dataset creation code](#). We will be making full use of these highly unusual and fun features in our visualizations, but let us first define what is meant by each feature as they are not all self-evident.

- Danceability: How suitable a track is for dancing. (0.0 -> 1.0).
- Energy: A perceptual measure of intensity and activity (0.0 -> 1.0).
- Speechiness: The presence of spoken words in a track. (0.0 -> 1.0).
- Acousticness: A confidence measure of whether a track is acoustic (0.0 -> 1.0).
- Instrumentalness: Predicts whether a track contains no vocals (0.0 -> 1.0).
- Liveness: Detects the presence of an audience in the recording (0.0 -> 1.0).
- Valence: Describes how positive the track is conveyed to be (0.0 -> 1.0).
- Loudness: The overall loudness of the track in decibels (-34.5 -> 1.51).

Let's inspect the shape of this data

```
[20]: np.shape(df)
```

```
[20]: (6513, 19)
```

Indeed, there are 6513 different rows, indicating 6513 tracks, and 19 different columns, uniquely indexed by the ID column. We would like to first check that there are actually 6513 distinct tracks,

```
[21]: unique_tracks = df['track_name'].nunique()
      print(f"Number of unique tracks: {unique_tracks}")
```

```
Number of unique tracks: 5351
```

As we might expect, there are fewer unique tracks than there are rows, suggesting the possibility for duplicate tracks. In practice this makes sense as some tracks share the same name but are different songs (for example “Afterglow” - by Ed Sheeran, and “Afterglow” - by Taylor Swift). We

might similarly expect the same, or a vastly similar song to be “remixed” and published by multiple artists. Lastly, some artists will even re-release some of their older songs, an example of this is taylor swift’s album “Fearless” - Taylor’s version.

Let’s see if there are any null values in the dataset.

```
[22]: df.isnull().sum()
```

```
[22]: id                0
      artist_names      0
      track_name        0
      source            0
      key               0
      mode              0
      time_signature    0
      danceability       0
      energy             0
      speechiness        0
      acousticness       0
      instrumentalness   0
      liveness           0
      valence            0
      loudness           0
      tempo              0
      duration_ms        0
      weeks_on_chart     0
      streams            0
      dtype: int64
```

We observe that there are no null values in the dataset which is nice to see!

Next we would like to start creating new dataframes from this raw dataframe so that it is more fit to be used for our purpose. To this end, we will need two separate dataframes, one that is indexed by artists, and one that is indexed by the tracks, as the one we have above already is what we are looking for, we will proceed by creating a dataframe that is indexed by the artists.

Our approach for doing this will be as follows: \* We will split artist\_names into artists, so that we have at most one artist per row. This will lead to some tracks being repeated, but that is fine. \* We will split the features into categorical features and continuous numerical columns (excluding weeks\_on\_chart and streams), and we will take the mean for the numerical features such as energy, on an artist to artist basis, and we will take the mode of the categorical features on an artist to artist basis. instead of averaging streams and weeks\_on\_charts, we think it is more suitable to aggregate and sum the streams. \* Finally we sort by the number of streams for each artist.

```
[28]: # Splitting artists and expanding the DataFrame to have one artist per row
      df['artist'] = df['artist_names'].str.split(',')
      df_expanded = df.explode('artist')

      # Drop the 'id' column and the original 'artist_names' column
      df_expanded.drop(['id', 'artist_names'], axis=1, inplace=True)
```

```

# Update numerical columns list, excluding 'weeks_on_chart' and 'streams' for
↳ special treatment
numerical_cols = ['danceability', 'energy', 'speechiness', 'acousticness',
                  'instrumentalness', 'liveness', 'valence', 'loudness',
                  'tempo', 'duration_ms']

# Define categorical columns
categorical_cols = ['key', 'mode', 'time_signature']

# Update aggregation functions
aggregations = {**{col: 'mean' for col in numerical_cols},
                **{col: lambda x: pd.Series.mode(x)[0] if not pd.Series.mode(x).
↳ empty else np.nan for col in categorical_cols},
                'weeks_on_chart': 'sum', # Summing weeks_on_chart
                'streams': 'sum'} # Summing streams

# Performing the aggregation
artist_df = df_expanded.groupby('artist', as_index=False).agg(aggregations)

# Rename columns 'weeks_on_chart' and 'streams' for clarity
artist_df.rename(columns={'weeks_on_chart': 'weeks_on_chart_total', 'streams':
↳ 'total_streams'}, inplace=True)

# Sort the DataFrame by 'total_streams' in descending order
artist_df_sorted = artist_df.sort_values(by='total_streams', ascending=False)

artist_df_sorted.head(10)

```

```

[28]:
      artist  danceability  energy  speechiness  acousticness \
172   Bad Bunny    0.749264  0.675874    0.123582    0.240423
1810  The Weeknd    0.612779  0.639500    0.097388    0.192148
509    Drake      0.706118  0.544405    0.194416    0.201880
1462  Post Malone  0.633874  0.637225    0.081846    0.261542
1757  Taylor Swift  0.596606  0.577000    0.062285    0.307204
535    Ed Sheeran  0.672726  0.610749    0.092308    0.323409
892   Justin Bieber  0.668372  0.587333    0.083179    0.276969
772    J Balvin    0.750067  0.729067    0.128871    0.126857
515    Dua Lipa    0.728800  0.722178    0.085451    0.064339
128   Ariana Grande  0.657109  0.614119    0.101528    0.233330

      instrumentalness  liveness  valence  loudness  tempo \
172          0.012356  0.172444  0.497427 -5.591346  122.768327
1810         0.007456  0.198514  0.389238 -6.732298  121.133731
509          0.013069  0.200462  0.366806 -7.848524  118.929795
1462         0.005981  0.159996  0.393936 -5.367523  122.025838
1757         0.004629  0.139067  0.401780 -7.631493  123.227473

```

535	0.000282	0.172138	0.534940	-6.035131	110.235393
892	0.004377	0.148803	0.550290	-6.618321	122.068397
772	0.008111	0.187484	0.647556	-4.684467	123.145311
515	0.000172	0.143736	0.647844	-5.124644	116.704222
128	0.000670	0.177766	0.472035	-6.133743	117.480604

	duration_ms	key	mode	time_signature	weeks_on_chart_total	\
172	214440.276730	C#/Db	Major	4 beats		2912
1810	220141.932692	C	Minor	4 beats		2243
509	226178.131004	C#/Db	Major	4 beats		2436
1462	194904.702703	C	Major	4 beats		2213
1757	237324.142857	C	Major	4 beats		1545
535	217245.000000	G#/Ab	Major	4 beats		2080
892	189779.294872	F#/Gb	Major	4 beats		1602
772	215137.966667	A	Minor	4 beats		1891
515	202467.600000	B	Minor	4 beats		1447
128	191435.039604	C	Major	4 beats		1308

	total_streams
172	31355337608
1810	21530493586
509	20961364679
1462	17775430798
1757	17106209239
535	16646366081
892	15666378558
772	14396444218
515	13890564424
128	12850455403

Unsurprisingly, the more popular artists have the most number of streams. It would now be interesting to take the top 50 most popular artists and find the top 5 artists for each “interesting” feature. Let us proceed.

```
[31]: # Isolating the top 50 most popular artists
top_50_artists = artist_df_sorted.head(50)

# Define "interesting" features
interesting_features = ['danceability', 'energy', 'speechiness', 'acousticness',
                       'instrumentalness', 'liveness', 'valence', 'loudness',
                       ↪ 'tempo']

# Initialize a dictionary to hold the top 5 artists for each feature
top_5_artists_per_feature = {}

# For each feature, find the top 5 artists
for feature in interesting_features:
```

```

top_5_artists = top_50_artists.nlargest(5, feature)[['artist', feature]]
top_5_artists_per_feature[feature] = top_5_artists

# Now, top_5_artists_per_feature contains the top 5 artists for each
↳ interesting feature
# Let's display the results
for feature, top_5_df in top_5_artists_per_feature.items():
    print(f"Top 5 artists for {feature}:")
    print(top_5_df)
    print("\n---\n")

```

Top 5 artists for danceability:

	artist	danceability
305	Cardi B	0.837542
414	DaBaby	0.817590
1340	Nicki Minaj	0.795565
1977	Young Thug	0.793333
904	KAROL G	0.781964

---

Top 5 artists for energy:

	artist	energy
415	Daddy Yankee	0.802250
1439	Peso Pluma	0.743395
118	Anuel AA	0.734823
1403	Ozuna	0.733974
1180	Marshmello	0.730393

---

Top 5 artists for speechiness:

	artist	speechiness
563	Eminem	0.267535
414	DaBaby	0.251990
944	Kendrick Lamar	0.236410
8	21 Savage	0.227059
1340	Nicki Minaj	0.210558

---

Top 5 artists for acousticness:

	artist	acousticness
213	Billie Eilish	0.674056
1027	Lewis Capaldi	0.632300
1596	Sam Smith	0.467740
1383	Olivia Rodrigo	0.455837
1947	XXXTENTACION	0.402349

---

Top 5 artists for instrumentalness:

	artist	instrumentalness
213	Billie Eilish	0.161094
954	Khalid	0.073047
363	Coldplay	0.042076
1947	XXXTENTACION	0.036406
8	21 Savage	0.021267

---

Top 5 artists for liveness:

	artist	liveness
563	Eminem	0.288039
415	Daddy Yankee	0.215920
363	Coldplay	0.214458
1853	Travis Scott	0.207632
509	Drake	0.200462

---

Top 5 artists for valence:

	artist	valence
1439	Peso Pluma	0.750070
1158	Maluma	0.700223
415	Daddy Yankee	0.698659
414	DaBaby	0.676256
515	Dua Lipa	0.647844

---

Top 5 artists for loudness:

	artist	loudness
415	Daddy Yankee	-3.861568
1506	Rauw Alejandro	-3.974327
118	Anuel AA	-4.389758
1158	Maluma	-4.452564
1403	Ozuna	-4.461092

---

Top 5 artists for tempo:

	artist	tempo
1042	Lil Nas X	140.948441
1383	Olivia Rodrigo	136.295677
643	Future	132.135346



```
881      Juice WRLD  131.735691
1439     Peso Pluma  131.702070
```

---

Let's visualize this data!

```
[32]: # Set up the matplotlib figure
fig, axes = plt.subplots(nrows=3, ncols=3, figsize=(20, 15),
    ↳ constrained_layout=True)
axes = axes.flatten()

# Iterate over each interesting feature and its corresponding top 5 DataFrame
for i, (feature, top_5_df) in enumerate(top_5_artists_per_feature.items()):
    sns.barplot(x=feature, y='artist', data=top_5_df, ax=axes[i])
    axes[i].set_title(f'Top 5 Artists for {feature.capitalize()}')
    axes[i].set_xlabel('Score')
    axes[i].set_ylabel('')

#plot
plt.tight_layout()
plt.show()
```

```
C:\Users\Christopher\AppData\Local\Temp\ipykernel_302612\1380144436.py:13:
UserWarning: The figure layout has changed to tight
plt.tight_layout()
```



Observing this data, we can kind of make sense of it, we would expect rappers like eminem to have the most "speechiness" in their songs on average. We would also expect the popular songs in the night club scene to appear on the "danceability" top 5.